

## Problem Definition :

User should be able to provide a Customer ID and Date, and program should be able to predict quantity

In [1]:

```
import pandas as pd
import numpy as np
```

In [2]:

```
data = pd.read_csv("online_retail.csv")
```

In [3]:

```
data.head()
```

Out[3]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T- LIGHT HOLDER	6	12/1/10 8:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	12/1/10 8:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/10 8:26	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/10 8:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/10 8:26	3.39	17850.0	United Kingdom

In [4]:

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   InvoiceNo       541909 non-null object
 1   StockCode      541909 non-null object
 2   Description     540455 non-null object
 3   Quantity       541909 non-null int64
 4   InvoiceDate     541909 non-null object
 5   UnitPrice      541909 non-null float64
 6   CustomerID     406829 non-null float64
 7   Country        541909 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 33.1+ MB
```

In [5]:

data.isna().sum()

Out[5]:

```
InvoiceNo      0
StockCode      0
Description    1454
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID    135080
Country        0
dtype: int64
```

In [6]:

```
# removing rows where description is null , customerid is null , unit price is less than
# all condition needs to be true
to_drop = data[(data['Description'].isnull()) & (data['CustomerID'].isnull()) &
               (data['UnitPrice'] < 1) & (data['Quantity'] < 0)].index

# Remove the filtered rows
data.drop(to_drop, inplace=True)
```

In [7]:

data.isna().sum()

Out[7]:

```
InvoiceNo      0
StockCode      0
Description     592
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID    134218
Country        0
dtype: int64
```

In [8]:

```
# forward filled all null CustomerID because
```

```
data['CustomerID'].fillna(method='ffill', inplace=True)
```

In [9]:

```
data.isna().sum()
```

Out[9]:

```
InvoiceNo      0
StockCode      0
Description    592
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID     0
Country        0
dtype: int64
```

In [10]:

```
# removing rows where quantity is in negative
# removing rows where quantity is greater than 2000 to remove outliers
```

```
data = data[(data['Quantity'] >= 0) & (data['Quantity'] <= 2000)]
```

In [11]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 531259 entries, 0 to 541908
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   InvoiceNo       531259 non-null object
 1   StockCode      531259 non-null object
 2   Description    530669 non-null object
 3   Quantity       531259 non-null int64
 4   InvoiceDate     531259 non-null object
 5   UnitPrice      531259 non-null float64
 6   CustomerID     531259 non-null float64
 7   Country        531259 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 36.5+ MB
```

In [12]:

```
neg_rows = data[data['Quantity'] <= 0]
neg_rows
```

Out[12]:

InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
-----------	-----------	-------------	----------	-------------	-----------	------------	---------

In [13]:

```
print(data.duplicated().sum())
data.drop_duplicates(inplace=True)
```

5231

In [14]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 526028 entries, 0 to 541908
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   InvoiceNo       526028 non-null object
 1   StockCode      526028 non-null object
 2   Description     525438 non-null object
 3   Quantity       526028 non-null int64
 4   InvoiceDate     526028 non-null object
 5   UnitPrice      526028 non-null float64
 6   CustomerID     526028 non-null float64
 7   Country        526028 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 36.1+ MB
```

In [15]:

```
# converted InvoiceDate to date time format and extracted Day, Month and Year.

data['InvoiceDate'] = pd.to_datetime(data['InvoiceDate'])
data['Day'] = data['InvoiceDate'].dt.day
data['Month'] = data['InvoiceDate'].dt.month
data['Year'] = data['InvoiceDate'].dt.year
```

In [16]:

```
# these columns are not required for prediction hence we are removing these columns.

data.drop(columns=['InvoiceNo', 'StockCode', 'Description', 'UnitPrice', 'Country', 'InvoiceDate'])
```

In [17]:

```
data.head(1)
```

Out[17]:

	Quantity	CustomerID	Day	Month	Year
0	6	17850.0	1	12	2010

In [18]:

```
data['CustomerID'] = data['CustomerID'].astype('int64')
data['Quantity'] = data['Quantity'].astype('int64')
```

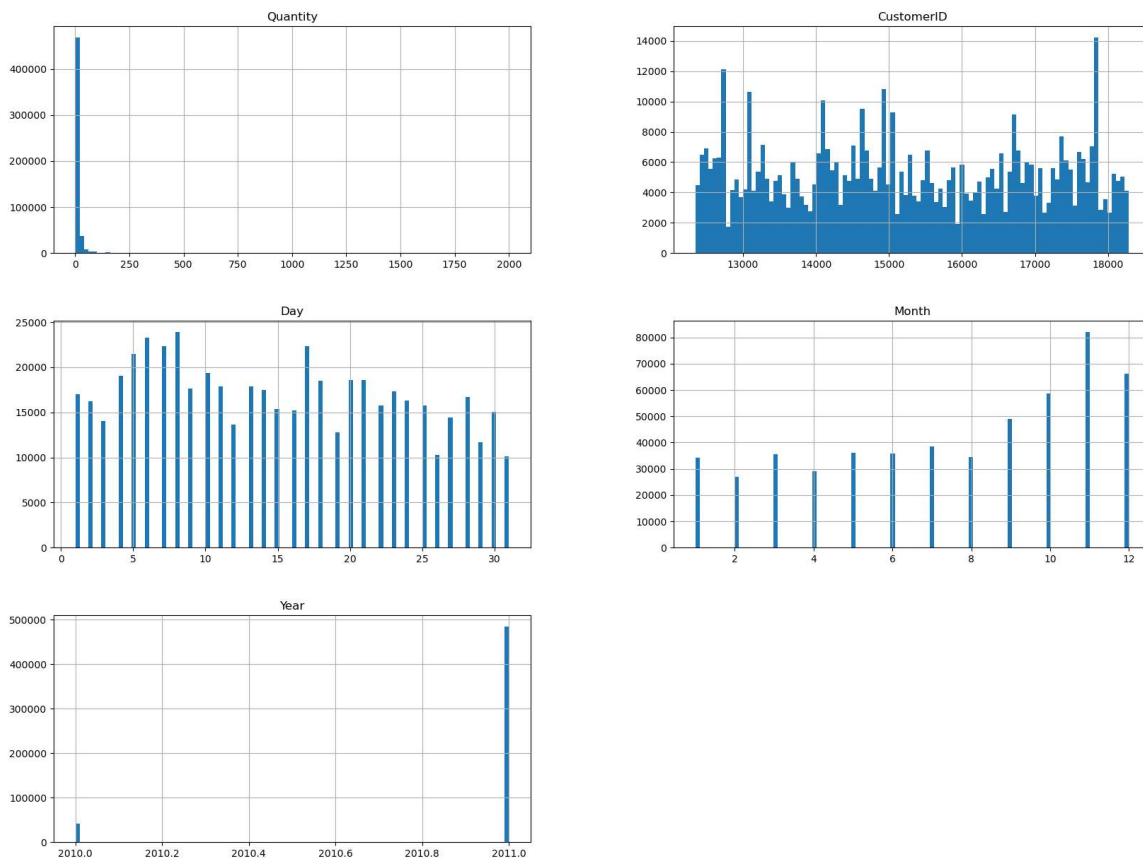
In [19]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 526028 entries, 0 to 541908
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   Quantity    526028 non-null  int64  
 1   CustomerID  526028 non-null  int64  
 2   Day          526028 non-null  int64  
 3   Month        526028 non-null  int64  
 4   Year         526028 non-null  int64  
dtypes: int64(5)
memory usage: 24.1 MB
```

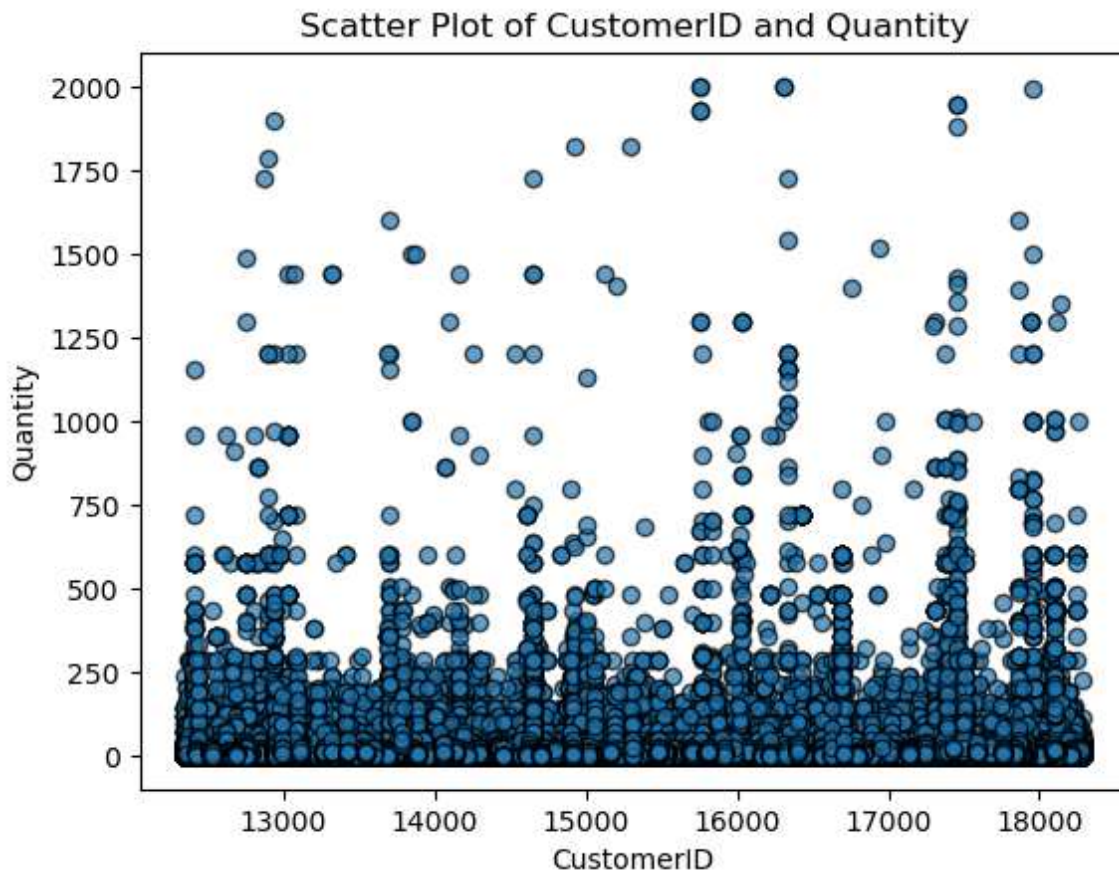
In [20]:

```
import matplotlib.pyplot as plt
data.hist(bins=100, figsize=(20,15))
plt.show()
```



In [21]:

```
plt.scatter(data['CustomerID'], data['Quantity'], edgecolor='black', alpha=0.7)  
plt.xlabel('CustomerID')  
plt.ylabel('Quantity')  
plt.title('Scatter Plot of CustomerID and Quantity')  
plt.show()
```



## Pivot Tables

In [22]:

```
pivot_table2 = pd.pivot_table(data, index='Day', columns='Year', values='Quantity', aggfun
pivot_table2.sort_values(by='All', ascending=False, inplace=True)
print(pivot_table2)
```

Year	2010	2011	All
Day			
All	355814.0	5050855.0	5406669
7	25324.0	215160.0	240484
5	16243.0	215497.0	231740
6	21775.0	194258.0	216033
20	15767.0	199456.0	215223
17	16947.0	195722.0	212669
11	NaN	211665.0	211665
14	20738.0	187885.0	208623
4	NaN	206013.0	206013
8	23049.0	175415.0	198464
21	15650.0	180715.0	196365
18	NaN	191519.0	191519
28	NaN	185337.0	185337
10	21031.0	161812.0	182843
9	20698.0	161086.0	181784
1	26906.0	152459.0	179365
15	18453.0	157725.0	176178
13	17701.0	153767.0	171468
24	NaN	168129.0	168129
16	27464.0	135904.0	163368
22	3212.0	157995.0	161207
2	28403.0	130329.0	158732
3	16430.0	138540.0	154970
19	3734.0	150583.0	154317
23	5752.0	146864.0	152616
25	NaN	149989.0	149989
12	10537.0	134570.0	145107
27	NaN	135642.0	135642
30	NaN	127676.0	127676
29	NaN	117130.0	117130
26	NaN	115819.0	115819
31	NaN	96194.0	96194

In [23]:

```
pivot_table3 = pd.pivot_table(data, index='Month', columns='Year', values='Quantity', agg
print(pivot_table3)
```

Year	2010	2011
Month		
1	NaN	311573.0
2	NaN	274776.0
3	NaN	378863.0
4	NaN	308200.0
5	NaN	391686.0
6	NaN	389633.0
7	NaN	399187.0
8	NaN	421386.0
9	NaN	574169.0
10	NaN	619173.0
11	NaN	748788.0
12	355814.0	233421.0

In [24]:

```
pivot_table_6 = data.pivot_table(index='CustomerID', columns='Year', values='Quantity', a
# Sorted pivot table high to low
pivot_table_6.sort_values(by='All', ascending=False, inplace=True)

print(pivot_table_6.head(10))
```

Year	2010	2011	All
CustomerID			
All	355814.0	5050855.0	5406669
14646	7983.0	188462.0	196445
14911	3331.0	81894.0	85225
12415	NaN	77812.0	77812
17450	786.0	70251.0	71037
18102	6695.0	61209.0	67904
13694	3860.0	62975.0	66835
17511	6426.0	59565.0	65991
14156	1606.0	59943.0	61549
14298	4976.0	53367.0	58343

## Model Implementation

In [25]:

```
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split
```

In [26]:

```
X = data[['CustomerID', 'Year', 'Month', 'Day']]
y = data['Quantity']
```



In [27]:

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(X)
X = pd.DataFrame(x_scaled, columns=X.columns)
```

In [28]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print(f"rows in train set : {len(X_train)}\nrows in test set : {len(X_test)}\n")
```

```
rows in train set : 420822
rows in test set : 105206
```

In [29]:

```
# regressor = LinearRegression()
# regressor.fit(X_train, y_train)

# regressor = RandomForestRegressor()
# regressor.fit(X_train, y_train)

regressor = DecisionTreeRegressor()
regressor.fit(X_train, y_train)
```

Out[29]:

```
DecisionTreeRegressor()
```

In [30]:

```
y_pred = regressor.predict(X_test)
```

In [31]:

```
print('R-squared score:', regressor.score(X_test, y_test))
```

```
R-squared score: 0.3267892862224361
```

In [32]:

```
from sklearn.metrics import mean_squared_error

# Calculating the MSE for the training data
y_train_pred = regressor.predict(X_train)
train_mse = mean_squared_error(y_train, y_train_pred)
print('Train MSE:', train_mse)
```

```
Train MSE: 489.944470249676
```

In [33]:

```
# Calculating the MSE for the test data
y_test_pred = regressor.predict(X_test)
test_mse = mean_squared_error(y_test, y_test_pred)
print('Test MSE:', test_mse)
```

Test MSE: 767.1758034959994

In [34]:

```
data = [[17850, 2010, 12, 1]] # CustomerID, Year, Month and Day

predicted_quantity = regressor.predict(data)

print('Predicted quantity:', predicted_quantity[0])
```

Predicted quantity: 2.6153846153846154

C:\Users\kajal\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but DecisionTreeRegressor was fitted with feature names  
warnings.warn(

In [ ]:

In [ ]: